

All the codes used in this analysis were gotten from tensorflow website, Kaggle notebook, Medium article and tutorials and Github repositories

1. Business Understanding

- **Objective:** Detect COVID-19 using medical imaging.
- **Evaluate** and compare the performance of various transfer learning models.

2. Data Understanding

2.1. Data Loading and Visualisation

```
In [1]: import os
import random
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.regularizers import l1_l2, l2
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.applications import MobileNetV2, DenseNet121
from tensorflow.keras.models import Model
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense, Dropout, BatchNormalisation
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ReduceLROnPlateau, EarlyStopping, LearningRateScheduler
from sklearn.metrics import classification_report, confusion_matrix, ConfusionMatrixDisplay
import seaborn as sns
import cv2

# Set random seeds for reproducibility
seed_value = 42
tf.random.set_seed(seed_value)
np.random.seed(seed_value)
random.seed(seed_value)
```

```
In [2]: data_dir = dataset_path = 'C:/Users/isacl/Downloads/Deep Learning project/COVID-19_
categories = ['COVID', 'Viral Pneumonia', 'Normal', 'Lung Opacity']
IMG_SIZE = 224
SAMPLE_SIZE = 12

valid_extensions = ['.jpg', '.jpeg', '.png']

def load_images_from_category(category, num_images):
    path = os.path.join(data_dir, category)
    images = []
    labels = []
```

```

label = categories.index(category)
count = 0

print(f"Loading images from: {path}")

for root, dirs, files in os.walk(path):
    for file in files:
        if count >= num_images:
            break
        if any(file.lower().endswith(ext) for ext in valid_extensions):
            img_path = os.path.join(root, file)
            print(f"Reading image: {img_path}")
            image = cv2.imread(img_path)
            if image is not None:
                image = cv2.resize(image, (IMG_SIZE, IMG_SIZE))
                image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB) # Convert to RGB
                image = image / 255.0 # Normalize
                images.append(image)
                labels.append(label)
                count += 1
            else:
                print(f"Failed to read image: {img_path}")

    return images, labels

def load_sample_images(sample_size_per_category):
    all_images = []
    all_labels = []

    for category in categories:
        images, labels = load_images_from_category(category, sample_size_per_category)
        all_images.extend(images)
        all_labels.extend(labels)

    return np.array(all_images), np.array(all_labels)

# Load a sample set of images
images, labels = load_sample_images(SAMPLE_SIZE // len(categories))
print(f"Loaded {len(images)} images.")

```

```

Loading images from: C:/Users/isacl/Downloads/Deep Learning project/COVID-19_Radiography_Dataset\COVID
Reading image: C:/Users/isacl/Downloads/Deep Learning project/COVID-19_Radiography_Dataset\COVID\images\COVID-1.png
Reading image: C:/Users/isacl/Downloads/Deep Learning project/COVID-19_Radiography_Dataset\COVID\images\COVID-10.png
Reading image: C:/Users/isacl/Downloads/Deep Learning project/COVID-19_Radiography_Dataset\COVID\images\COVID-100.png
Loading images from: C:/Users/isacl/Downloads/Deep Learning project/COVID-19_Radiography_Dataset\Viral Pneumonia
Reading image: C:/Users/isacl/Downloads/Deep Learning project/COVID-19_Radiography_Dataset\Viral Pneumonia\images\Viral Pneumonia-1.png
Reading image: C:/Users/isacl/Downloads/Deep Learning project/COVID-19_Radiography_Dataset\Viral Pneumonia\images\Viral Pneumonia-10.png
Reading image: C:/Users/isacl/Downloads/Deep Learning project/COVID-19_Radiography_Dataset\Viral Pneumonia\images\Viral Pneumonia-100.png
Loading images from: C:/Users/isacl/Downloads/Deep Learning project/COVID-19_Radiography_Dataset\Normal
Reading image: C:/Users/isacl/Downloads/Deep Learning project/COVID-19_Radiography_Dataset\Normal\images\Normal-1.png
Reading image: C:/Users/isacl/Downloads/Deep Learning project/COVID-19_Radiography_Dataset\Normal\images\Normal-10.png
Reading image: C:/Users/isacl/Downloads/Deep Learning project/COVID-19_Radiography_Dataset\Normal\images\Normal-100.png
Loading images from: C:/Users/isacl/Downloads/Deep Learning project/COVID-19_Radiography_Dataset\Lung Opacity
Reading image: C:/Users/isacl/Downloads/Deep Learning project/COVID-19_Radiography_Dataset\Lung Opacity\images\Lung Opacity-1.png
Reading image: C:/Users/isacl/Downloads/Deep Learning project/COVID-19_Radiography_Dataset\Lung Opacity\images\Lung Opacity-10.png
Reading image: C:/Users/isacl/Downloads/Deep Learning project/COVID-19_Radiography_Dataset\Lung Opacity\images\Lung Opacity-100.png
Loaded 12 images.

```

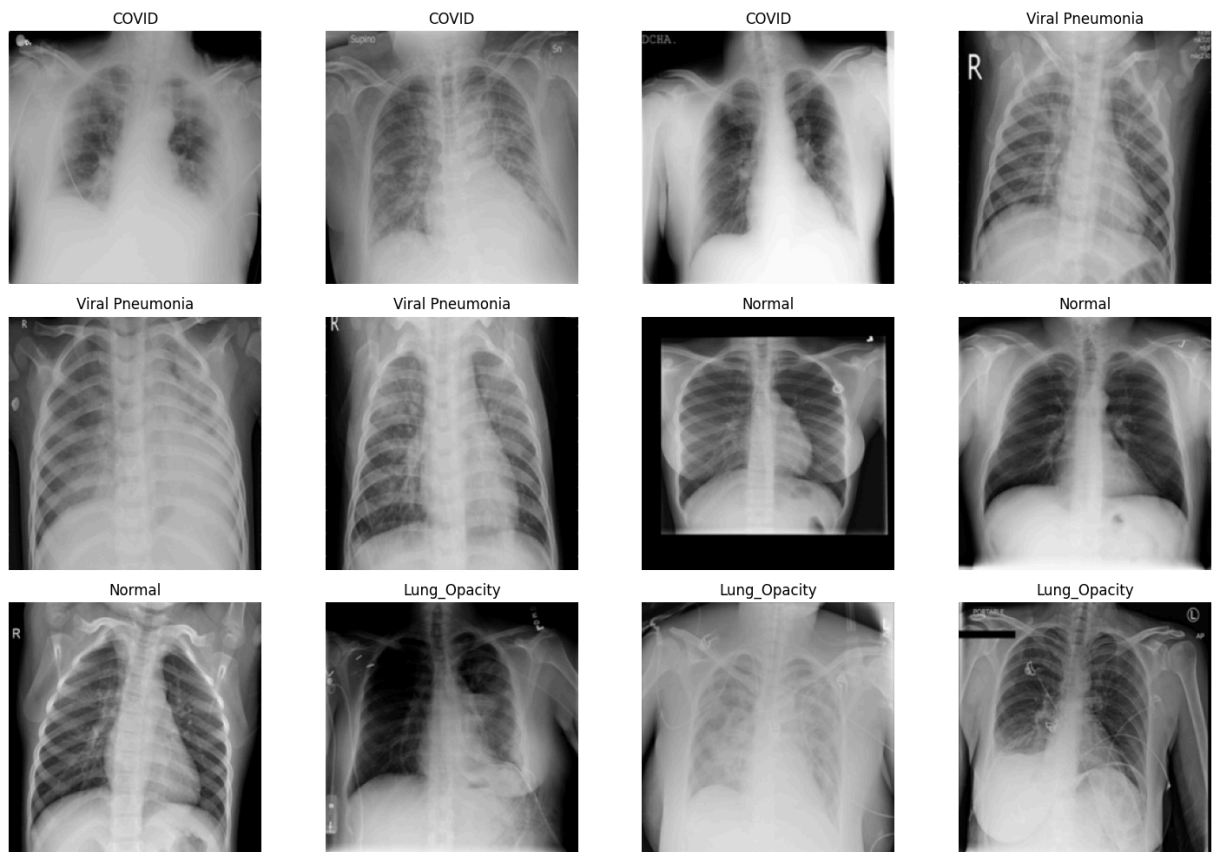
```

In [3]: def visualize_sample_images(images, labels, categories):
        if len(images) == 0:
            print("No images to display.")
            return

        plt.figure(figsize=(15, 10))
        for i in range(len(images)):
            plt.subplot(3, 4, i + 1)
            plt.imshow(images[i])
            plt.title(categories[labels[i]], fontsize=12)
            plt.axis('off')
        plt.tight_layout()
        plt.show()

        visualize_sample_images(images, labels, categories)

```



3. Data Preprocessing

3.1. Creating Data Generator

```
In [4]: # Create data generator with augmentation
datagen = ImageDataGenerator(
    rescale=1./255,
    validation_split=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True
)

train_generator = datagen.flow_from_directory(
    dataset_path,
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical',
    subset='training'
)

validation_generator = datagen.flow_from_directory(
    dataset_path,
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical',
    subset='validation'
```

```

)

# Verify dataset sizes and batch sizes
print("Number of training samples:", train_generator.samples)
print("Number of validation samples:", validation_generator.samples)
print("Training batch size:", train_generator.batch_size)
print("Validation batch size:", validation_generator.batch_size)

# Calculate steps per epoch
steps_per_epoch_train = train_generator.samples // train_generator.batch_size
steps_per_epoch_val = validation_generator.samples // validation_generator.batch_size

print(f"Steps per epoch (train): {steps_per_epoch_train}")
print(f"Steps per epoch (val): {steps_per_epoch_val}")

# Use .repeat() to ensure the dataset generates enough batches
train_dataset = tf.data.Dataset.from_generator(
    lambda: (x for x in train_generator),
    output_signature=(
        tf.TensorSpec(shape=(None, 224, 224, 3), dtype=tf.float32),
        tf.TensorSpec(shape=(None, 4), dtype=tf.float32)
    )
).repeat().prefetch(tf.data.AUTOTUNE)

validation_dataset = tf.data.Dataset.from_generator(
    lambda: (x for x in validation_generator),
    output_signature=(
        tf.TensorSpec(shape=(None, 224, 224, 3), dtype=tf.float32),
        tf.TensorSpec(shape=(None, 4), dtype=tf.float32)
    )
).repeat().prefetch(tf.data.AUTOTUNE)

```

Found 33866 images belonging to 4 classes.

Found 8464 images belonging to 4 classes.

Number of training samples: 33866

Number of validation samples: 8464

Training batch size: 32

Validation batch size: 32

Steps per epoch (train): 1058

Steps per epoch (val): 264

4. Model Training and Evaluation

4.1. MobileNetV2 with L1 and L2 Regularization and Learning Rate Scheduling

```

In [5]: def create_mobilenetv2_model():
        base_model = MobileNetV2(input_shape=(224, 224, 3), include_top=False, weights=

        # Unfreeze the last 50 layers
        for layer in base_model.layers[:-50]:
            layer.trainable = False
        for layer in base_model.layers[-50:]:
            layer.trainable = True

```

```

model = Sequential([
    base_model,
    Flatten(),
    BatchNormalization(),
    Dense(128, activation='relu', kernel_regularizer=l1_l2(l1=0.001, l2=0.001)),
    Dropout(0.5),
    Dense(4, activation='softmax', kernel_regularizer=l1_l2(l1=0.001, l2=0.001))
])
optimizer = Adam(learning_rate=0.0001) # Adjusted Learning rate
model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['a
return model

mobilenetv2_model = create_mobilenetv2_model()

# Learning rate schedule
def scheduler(epoch, lr):
    if epoch < 10:
        return float(lr)
    else:
        return float(lr * tf.math.exp(-0.1))

lr_scheduler = LearningRateScheduler(scheduler)

early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights
mobilenetv2_checkpoint = ModelCheckpoint('best_mobilenetv2_model.keras', monitor='v

history_mobilenetv2 = mobilenetv2_model.fit(
    train_dataset,
    validation_data=validation_dataset,
    epochs=20,
    steps_per_epoch=steps_per_epoch_train,
    validation_steps=steps_per_epoch_val,
    callbacks=[early_stopping, mobilenetv2_checkpoint, lr_scheduler]
)

```

Epoch 1/20
1058/1058 ————— **1061s** 988ms/step - accuracy: 0.6091 - loss: 28.2470 - val_accuracy: 0.6838 - val_loss: 6.3597 - learning_rate: 1.0000e-04

Epoch 2/20
1058/1058 ————— **770s** 728ms/step - accuracy: 0.7335 - loss: 4.0839 - val_accuracy: 0.7939 - val_loss: 2.1837 - learning_rate: 1.0000e-04

Epoch 3/20
1058/1058 ————— **772s** 730ms/step - accuracy: 0.7559 - loss: 2.1561 - val_accuracy: 0.8744 - val_loss: 1.4531 - learning_rate: 1.0000e-04

Epoch 4/20
1058/1058 ————— **785s** 742ms/step - accuracy: 0.7703 - loss: 1.6094 - val_accuracy: 0.8987 - val_loss: 1.0936 - learning_rate: 1.0000e-04

Epoch 5/20
1058/1058 ————— **777s** 735ms/step - accuracy: 0.7839 - loss: 1.3422 - val_accuracy: 0.8255 - val_loss: 1.1396 - learning_rate: 1.0000e-04

Epoch 6/20
1058/1058 ————— **737s** 697ms/step - accuracy: 0.7926 - loss: 1.1764 - val_accuracy: 0.8021 - val_loss: 1.0337 - learning_rate: 1.0000e-04

Epoch 7/20
1058/1058 ————— **750s** 710ms/step - accuracy: 0.8055 - loss: 1.0777 - val_accuracy: 0.8987 - val_loss: 0.8451 - learning_rate: 1.0000e-04

Epoch 8/20
1058/1058 ————— **684s** 647ms/step - accuracy: 0.8101 - loss: 1.0467 - val_accuracy: 0.9105 - val_loss: 0.7206 - learning_rate: 1.0000e-04

Epoch 9/20
1058/1058 ————— **655s** 619ms/step - accuracy: 0.8078 - loss: 0.9847 - val_accuracy: 0.8898 - val_loss: 0.7474 - learning_rate: 1.0000e-04

Epoch 10/20
1058/1058 ————— **671s** 635ms/step - accuracy: 0.8197 - loss: 0.9168 - val_accuracy: 0.9097 - val_loss: 0.7131 - learning_rate: 1.0000e-04

Epoch 11/20
1058/1058 ————— **697s** 659ms/step - accuracy: 0.8272 - loss: 0.8660 - val_accuracy: 0.9126 - val_loss: 0.6741 - learning_rate: 9.0484e-05

Epoch 12/20
1058/1058 ————— **702s** 664ms/step - accuracy: 0.8323 - loss: 0.8251 - val_accuracy: 0.9275 - val_loss: 0.5815 - learning_rate: 8.1873e-05

Epoch 13/20
1058/1058 ————— **676s** 640ms/step - accuracy: 0.8410 - loss: 0.7655 - val_accuracy: 0.9210 - val_loss: 0.5771 - learning_rate: 7.4082e-05

Epoch 14/20
1058/1058 ————— **669s** 633ms/step - accuracy: 0.8410 - loss: 0.7345 - val_accuracy: 0.9280 - val_loss: 0.5173 - learning_rate: 6.7032e-05

Epoch 15/20
1058/1058 ————— **664s** 627ms/step - accuracy: 0.8492 - loss: 0.7005 - val_accuracy: 0.9181 - val_loss: 0.5181 - learning_rate: 6.0653e-05

Epoch 16/20
1058/1058 ————— **673s** 637ms/step - accuracy: 0.8560 - loss: 0.6577 - val_accuracy: 0.9348 - val_loss: 0.4607 - learning_rate: 5.4881e-05

Epoch 17/20
1058/1058 ————— **679s** 642ms/step - accuracy: 0.8598 - loss: 0.6321 - val_accuracy: 0.9320 - val_loss: 0.4603 - learning_rate: 4.9659e-05

Epoch 18/20
1058/1058 ————— **676s** 639ms/step - accuracy: 0.8654 - loss: 0.6150 - val_accuracy: 0.9286 - val_loss: 0.4699 - learning_rate: 4.4933e-05

Epoch 19/20
1058/1058 ————— **674s** 637ms/step - accuracy: 0.8678 - loss: 0.5929 - val_accuracy: 0.9286 - val_loss: 0.4699 - learning_rate: 4.4933e-05

al_accuracy: 0.9298 - val_loss: 0.4412 - learning_rate: 4.0657e-05

Epoch 20/20

1058/1058 ————— 663s 627ms/step - accuracy: 0.8725 - loss: 0.5672 - v

al_accuracy: 0.9272 - val_loss: 0.4636 - learning_rate: 3.6788e-05


```
In [6]: # Function to plot training history
def plot_training_history(history, title):
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']
    loss = history.history['loss']
    val_loss = history.history['val_loss']
    epochs = range(1, len(acc) + 1)

    plt.figure(figsize=(14, 5))
    plt.subplot(1, 2, 1)
    plt.plot(epochs, acc, 'b', label='Training accuracy')
    plt.plot(epochs, val_acc, 'r', label='Validation accuracy')
    plt.title(f'Training and validation accuracy: {title}')
    plt.legend()

    plt.subplot(1, 2, 2)
    plt.plot(epochs, loss, 'b', label='Training loss')
    plt.plot(epochs, val_loss, 'r', label='Validation loss')
    plt.title(f'Training and validation loss: {title}')
    plt.legend()

    plt.show()

# Plot history for MobileNetV2
plot_training_history(history_mobilenetv2, 'MobileNetV2')
```



```
In [7]: # Function to evaluate the model
def evaluate_model(model, validation_generator, steps):
    y_true = []
    y_pred = []

    for i in range(steps):
        x_val, y_val = next(validation_generator)
        preds = model.predict(x_val)
        y_true.extend(np.argmax(y_val, axis=1))
        y_pred.extend(np.argmax(preds, axis=1))
```



```
# Calculate confusion matrix
conf_matrix = confusion_matrix(y_true, y_pred)

# Plot confusion matrix heatmap
plt.figure(figsize=(10, 8))
disp = ConfusionMatrixDisplay(confusion_matrix=conf_matrix, display_labels=validation_generator.classes_)
disp.plot(cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.show()

# Print classification report
report = classification_report(y_true, y_pred, target_names=validation_generator.classes_)
print(report)

# Calculate and print metrics
accuracy = np.trace(conf_matrix) / np.sum(conf_matrix)
precision = np.diag(conf_matrix) / np.sum(conf_matrix, axis=0)
recall = np.diag(conf_matrix) / np.sum(conf_matrix, axis=1)
specificity = (np.sum(conf_matrix) - np.sum(conf_matrix, axis=0) - np.sum(conf_matrix, axis=1)) / (np.sum(conf_matrix) - np.sum(conf_matrix, axis=0) - np.sum(conf_matrix, axis=1))
f1_scores = 2 * (precision * recall) / (precision + recall)

print(f'Accuracy: {accuracy:.4f}')
print(f'Precision: {precision}')
print(f'Recall: {recall}')
print(f'Specificity: {specificity}')
print(f'F1-Scores: {f1_scores}')

# Evaluate MobileNetV2 model
evaluate_model(mobilenetv2_model, validation_generator, steps_per_epoch_val)
```

1/1	2s	2s/step
1/1	0s	250ms/step
1/1	0s	247ms/step
1/1	0s	250ms/step
1/1	0s	248ms/step
1/1	0s	250ms/step
1/1	0s	245ms/step
1/1	0s	240ms/step
1/1	0s	253ms/step
1/1	0s	247ms/step
1/1	0s	240ms/step
1/1	0s	237ms/step
1/1	0s	247ms/step
1/1	0s	247ms/step
1/1	0s	240ms/step
1/1	0s	252ms/step
1/1	0s	237ms/step
1/1	0s	250ms/step
1/1	0s	253ms/step
1/1	0s	240ms/step
1/1	0s	246ms/step
1/1	0s	245ms/step
1/1	0s	240ms/step
1/1	0s	243ms/step
1/1	0s	314ms/step
1/1	0s	352ms/step
1/1	0s	313ms/step
1/1	0s	255ms/step
1/1	0s	249ms/step
1/1	0s	240ms/step
1/1	0s	341ms/step
1/1	0s	307ms/step
1/1	0s	288ms/step
1/1	0s	285ms/step
1/1	0s	279ms/step
1/1	0s	281ms/step
1/1	0s	257ms/step
1/1	0s	246ms/step
1/1	0s	272ms/step
1/1	0s	257ms/step
1/1	0s	258ms/step
1/1	0s	255ms/step
1/1	0s	262ms/step
1/1	0s	243ms/step
1/1	0s	290ms/step
1/1	0s	288ms/step
1/1	0s	312ms/step
1/1	0s	255ms/step
1/1	0s	281ms/step
1/1	0s	256ms/step
1/1	0s	310ms/step
1/1	0s	287ms/step
1/1	0s	326ms/step
1/1	0s	256ms/step
1/1	0s	246ms/step
1/1	0s	247ms/step


1/1	0s	241ms/step
1/1	0s	245ms/step
1/1	0s	246ms/step
1/1	0s	240ms/step
1/1	0s	247ms/step
1/1	0s	247ms/step
1/1	0s	244ms/step
1/1	0s	242ms/step
1/1	0s	245ms/step
1/1	0s	240ms/step
1/1	0s	237ms/step
1/1	0s	240ms/step
1/1	0s	238ms/step
1/1	0s	244ms/step
1/1	0s	245ms/step
1/1	0s	237ms/step
1/1	0s	329ms/step
1/1	0s	254ms/step
1/1	0s	323ms/step
1/1	0s	282ms/step
1/1	0s	247ms/step
1/1	0s	233ms/step
1/1	0s	249ms/step
1/1	0s	245ms/step
1/1	0s	240ms/step
1/1	0s	240ms/step
1/1	0s	250ms/step
1/1	0s	280ms/step
1/1	0s	283ms/step
1/1	0s	246ms/step
1/1	0s	240ms/step
1/1	0s	266ms/step
1/1	0s	244ms/step
1/1	0s	253ms/step
1/1	0s	260ms/step
1/1	0s	282ms/step
1/1	0s	275ms/step
1/1	0s	282ms/step
1/1	0s	278ms/step
1/1	0s	280ms/step
1/1	0s	266ms/step
1/1	0s	270ms/step
1/1	0s	290ms/step
1/1	0s	273ms/step
1/1	0s	243ms/step
1/1	0s	300ms/step
1/1	0s	251ms/step
1/1	0s	247ms/step
1/1	0s	233ms/step
1/1	0s	244ms/step
1/1	0s	244ms/step
1/1	0s	250ms/step
1/1	0s	248ms/step
1/1	0s	246ms/step
1/1	0s	235ms/step
1/1	0s	246ms/step

1/1	0s	245ms/step
1/1	0s	238ms/step
1/1	0s	252ms/step
1/1	0s	243ms/step
1/1	0s	244ms/step
1/1	0s	240ms/step
1/1	0s	246ms/step
1/1	0s	324ms/step
1/1	0s	296ms/step
1/1	0s	307ms/step
1/1	0s	339ms/step
1/1	0s	335ms/step
1/1	0s	245ms/step
1/1	0s	242ms/step
1/1	0s	281ms/step
1/1	0s	244ms/step
1/1	0s	239ms/step
1/1	0s	246ms/step
1/1	0s	236ms/step
1/1	0s	234ms/step
1/1	0s	245ms/step
1/1	0s	259ms/step
1/1	0s	244ms/step
1/1	0s	252ms/step
1/1	0s	255ms/step
1/1	0s	256ms/step
1/1	0s	260ms/step
1/1	0s	243ms/step
1/1	0s	243ms/step
1/1	0s	251ms/step
1/1	0s	255ms/step
1/1	0s	250ms/step
1/1	0s	255ms/step
1/1	0s	250ms/step
1/1	0s	246ms/step
1/1	0s	259ms/step
1/1	0s	258ms/step
1/1	0s	255ms/step
1/1	0s	254ms/step
1/1	0s	255ms/step
1/1	0s	254ms/step
1/1	0s	251ms/step
1/1	0s	252ms/step
1/1	0s	246ms/step
1/1	0s	261ms/step
1/1	0s	250ms/step
1/1	0s	254ms/step
1/1	0s	250ms/step
1/1	0s	262ms/step
1/1	0s	243ms/step
1/1	0s	265ms/step
1/1	0s	252ms/step
1/1	0s	247ms/step
1/1	0s	257ms/step
1/1	0s	255ms/step
1/1	0s	255ms/step

1/1	0s	243ms/step
1/1	0s	253ms/step
1/1	0s	266ms/step
1/1	0s	271ms/step
1/1	0s	254ms/step
1/1	0s	257ms/step
1/1	0s	250ms/step
1/1	0s	261ms/step
1/1	0s	257ms/step
1/1	0s	254ms/step
1/1	0s	237ms/step
1/1	0s	258ms/step
1/1	0s	265ms/step
1/1	0s	255ms/step
1/1	0s	259ms/step
1/1	0s	274ms/step
1/1	0s	264ms/step
1/1	0s	250ms/step
1/1	0s	260ms/step
1/1	0s	250ms/step
1/1	0s	260ms/step
1/1	0s	247ms/step
1/1	0s	250ms/step
1/1	0s	250ms/step
1/1	0s	261ms/step
1/1	0s	260ms/step
1/1	0s	256ms/step
1/1	0s	261ms/step
1/1	0s	250ms/step
1/1	0s	246ms/step
1/1	0s	252ms/step
1/1	0s	255ms/step
1/1	0s	254ms/step
1/1	0s	253ms/step
1/1	0s	256ms/step
1/1	0s	252ms/step
1/1	0s	245ms/step
1/1	0s	264ms/step
1/1	0s	261ms/step
1/1	0s	250ms/step
1/1	2s	2s/step
1/1	0s	260ms/step
1/1	0s	255ms/step
1/1	0s	250ms/step
1/1	0s	250ms/step
1/1	0s	260ms/step
1/1	0s	242ms/step
1/1	0s	254ms/step
1/1	0s	255ms/step
1/1	0s	256ms/step
1/1	0s	247ms/step
1/1	0s	250ms/step
1/1	0s	255ms/step
1/1	0s	257ms/step
1/1	0s	246ms/step
1/1	0s	264ms/step

1/1	—————	0s	266ms/step
1/1	—————	0s	261ms/step
1/1	—————	0s	259ms/step
1/1	—————	0s	261ms/step
1/1	—————	0s	261ms/step
1/1	—————	0s	251ms/step
1/1	—————	0s	343ms/step
1/1	—————	0s	243ms/step
1/1	—————	0s	263ms/step
1/1	—————	0s	367ms/step
1/1	—————	0s	290ms/step
1/1	—————	0s	258ms/step
1/1	—————	0s	255ms/step
1/1	—————	0s	265ms/step
1/1	—————	0s	260ms/step
1/1	—————	0s	250ms/step
1/1	—————	0s	249ms/step
1/1	—————	0s	255ms/step
1/1	—————	0s	260ms/step
1/1	—————	0s	243ms/step
1/1	—————	0s	245ms/step
1/1	—————	0s	242ms/step
1/1	—————	0s	242ms/step
1/1	—————	0s	236ms/step
1/1	—————	0s	240ms/step
1/1	—————	0s	245ms/step
1/1	—————	0s	249ms/step
1/1	—————	0s	240ms/step
1/1	—————	0s	241ms/step
1/1	—————	0s	250ms/step
1/1	—————	0s	246ms/step
1/1	—————	0s	238ms/step
1/1	—————	0s	257ms/step
1/1	—————	0s	240ms/step
1/1	—————	0s	250ms/step
1/1	—————	0s	245ms/step
1/1	—————	0s	244ms/step
1/1	—————	0s	239ms/step
1/1	—————	0s	250ms/step
1/1	—————	0s	247ms/step

<Figure size 1000x800 with 0 Axes>



	precision	recall	f1-score	support
COVID	0.97	0.91	0.94	1426
Lung_Opacity	0.92	0.91	0.92	2411
Normal	0.94	0.95	0.94	4050
Viral Pneumonia	0.89	0.98	0.93	545
accuracy			0.93	8432
macro avg	0.93	0.94	0.93	8432
weighted avg	0.93	0.93	0.93	8432

Accuracy: 0.9343

Precision: [0.96569724 0.92352941 0.93740867 0.88595041]

Recall: [0.90813464 0.91165491 0.95037037 0.98348624]

Specificity: [0.9934342 0.96977246 0.94135098 0.99125143]

F1-Scores: [0.9360318 0.91755375 0.94384502 0.93217391]

4.2 DenseNet121 Model

```
In [26]: # Load DenseNet121 model pre-trained on ImageNet
def create_densenet_model():
    base_model = DenseNet121(input_shape=(224, 224, 3), include_top=False, weights=

    # Unfreeze the last few layers for fine-tuning
    for layer in base_model.layers[:-50]:
        layer.trainable = False
    for layer in base_model.layers[-50:]:
        layer.trainable = True
```

```

model = Sequential([
    base_model,
    Flatten(),
    BatchNormalization(),
    Dense(128, activation='relu', kernel_regularizer=l2(0.001)),
    Dropout(0.5),
    Dense(4, activation='softmax', kernel_regularizer=l2(0.001))
])

optimizer = Adam(learning_rate=0.0001)
model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['a
return model

densenet_model = create_densenet_model()

# Callbacks for early stopping and saving the best model
early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights
model_checkpoint = ModelCheckpoint('best_densenet_model.keras', monitor='val_loss',

# Learning rate scheduler
def scheduler(epoch, lr):
    return float(lr * tf.math.exp(-0.1))

lr_scheduler = LearningRateScheduler(scheduler)

# Training the DenseNet model
history_densenet = densenet_model.fit(
    train_dataset,
    validation_data=validation_dataset,
    epochs=20,
    steps_per_epoch=steps_per_epoch_train,
    validation_steps=steps_per_epoch_val,
    callbacks=[early_stopping, model_checkpoint, lr_scheduler]
)

```


Epoch 1/20
1058/1058 ————— **1337s** 1s/step - accuracy: 0.6301 - loss: 1.4840 - val
_accuracy: 0.8647 - val_loss: 0.6600 - learning_rate: 9.0484e-05

Epoch 2/20
1058/1058 ————— **1266s** 1s/step - accuracy: 0.7337 - loss: 0.9238 - val
_accuracy: 0.8678 - val_loss: 0.6001 - learning_rate: 8.1873e-05

Epoch 3/20
1058/1058 ————— **1291s** 1s/step - accuracy: 0.7520 - loss: 0.8599 - val
_accuracy: 0.8907 - val_loss: 0.5240 - learning_rate: 7.4082e-05

Epoch 4/20
1058/1058 ————— **1474s** 1s/step - accuracy: 0.7706 - loss: 0.7938 - val
_accuracy: 0.8987 - val_loss: 0.4983 - learning_rate: 6.7032e-05

Epoch 5/20
1058/1058 ————— **1262s** 1s/step - accuracy: 0.7794 - loss: 0.7612 - val
_accuracy: 0.8982 - val_loss: 0.4833 - learning_rate: 6.0653e-05

Epoch 6/20
1058/1058 ————— **1248s** 1s/step - accuracy: 0.7844 - loss: 0.7365 - val
_accuracy: 0.8991 - val_loss: 0.4709 - learning_rate: 5.4881e-05

Epoch 7/20
1058/1058 ————— **1248s** 1s/step - accuracy: 0.7926 - loss: 0.7162 - val
_accuracy: 0.9012 - val_loss: 0.4644 - learning_rate: 4.9659e-05

Epoch 8/20
1058/1058 ————— **1237s** 1s/step - accuracy: 0.8028 - loss: 0.6799 - val
_accuracy: 0.9025 - val_loss: 0.4668 - learning_rate: 4.4933e-05

Epoch 9/20
1058/1058 ————— **1238s** 1s/step - accuracy: 0.8059 - loss: 0.6671 - val
_accuracy: 0.9099 - val_loss: 0.4238 - learning_rate: 4.0657e-05

Epoch 10/20
1058/1058 ————— **1249s** 1s/step - accuracy: 0.8093 - loss: 0.6414 - val
_accuracy: 0.9166 - val_loss: 0.4201 - learning_rate: 3.6788e-05

Epoch 11/20
1058/1058 ————— **1238s** 1s/step - accuracy: 0.8116 - loss: 0.6271 - val
_accuracy: 0.9071 - val_loss: 0.4183 - learning_rate: 3.3287e-05

Epoch 12/20
1058/1058 ————— **1249s** 1s/step - accuracy: 0.8178 - loss: 0.6069 - val
_accuracy: 0.9125 - val_loss: 0.4290 - learning_rate: 3.0119e-05

Epoch 13/20
1058/1058 ————— **1285s** 1s/step - accuracy: 0.8198 - loss: 0.5981 - val
_accuracy: 0.9171 - val_loss: 0.3965 - learning_rate: 2.7253e-05

Epoch 14/20
1058/1058 ————— **1283s** 1s/step - accuracy: 0.8234 - loss: 0.5832 - val
_accuracy: 0.9160 - val_loss: 0.4155 - learning_rate: 2.4660e-05

Epoch 15/20
1058/1058 ————— **1289s** 1s/step - accuracy: 0.8286 - loss: 0.5707 - val
_accuracy: 0.9177 - val_loss: 0.3959 - learning_rate: 2.2313e-05

Epoch 16/20
1058/1058 ————— **1301s** 1s/step - accuracy: 0.8291 - loss: 0.5645 - val
_accuracy: 0.9176 - val_loss: 0.3835 - learning_rate: 2.0190e-05

Epoch 17/20
1058/1058 ————— **1258s** 1s/step - accuracy: 0.8296 - loss: 0.5567 - val
_accuracy: 0.9249 - val_loss: 0.3941 - learning_rate: 1.8268e-05

Epoch 18/20
1058/1058 ————— **1243s** 1s/step - accuracy: 0.8342 - loss: 0.5390 - val
_accuracy: 0.9233 - val_loss: 0.3622 - learning_rate: 1.6530e-05

Epoch 19/20
1058/1058 ————— **1250s** 1s/step - accuracy: 0.8348 - loss: 0.5374 - val

_accuracy: 0.9227 - val_loss: 0.3721 - learning_rate: 1.4957e-05

Epoch 20/20

1058/1058 ————— **1292s** 1s/step - accuracy: 0.8387 - loss: 0.5306 - val
_accuracy: 0.9224 - val_loss: 0.3686 - learning_rate: 1.3534e-05

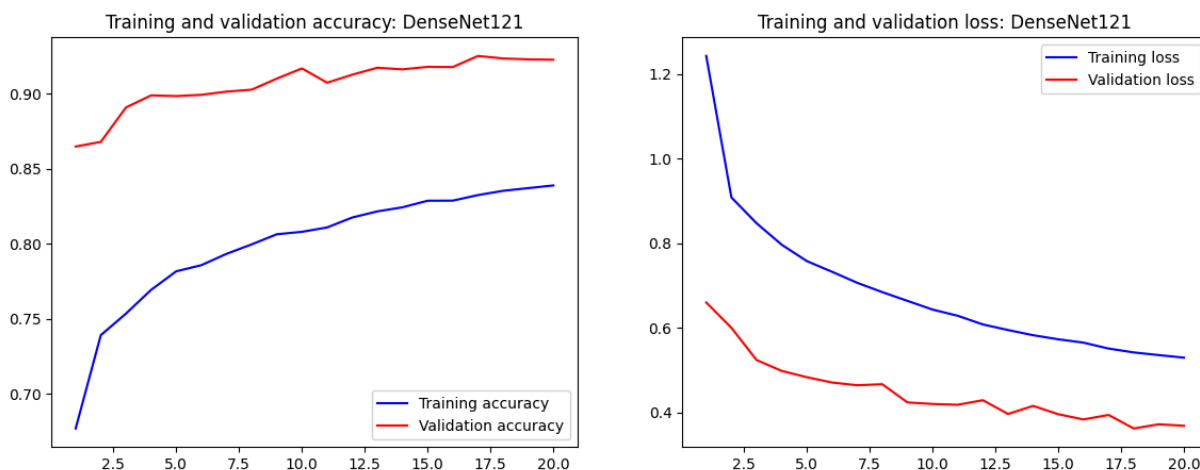
```
In [29]: # Function to plot training history
def plot_training_history(history, title):
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']
    loss = history.history['loss']
    val_loss = history.history['val_loss']
    epochs = range(1, len(acc) + 1)

    plt.figure(figsize=(14, 5))
    plt.subplot(1, 2, 1)
    plt.plot(epochs, acc, 'b', label='Training accuracy')
    plt.plot(epochs, val_acc, 'r', label='Validation accuracy')
    plt.title(f'Training and validation accuracy: {title}')
    plt.legend()

    plt.subplot(1, 2, 2)
    plt.plot(epochs, loss, 'b', label='Training loss')
    plt.plot(epochs, val_loss, 'r', label='Validation loss')
    plt.title(f'Training and validation loss: {title}')
    plt.legend()

    plt.show()

# Plot history for DenseNet121
plot_training_history(history_densenet, 'DenseNet121')
```



```
In [31]: # Function to evaluate the model
def evaluate_model(model, validation_generator, steps):
    y_true = []
    y_pred = []

    for i in range(steps):
        x_val, y_val = next(validation_generator)
        preds = model.predict(x_val)
        y_true.extend(np.argmax(y_val, axis=1))
        y_pred.extend(np.argmax(preds, axis=1))
```

```
# Calculate confusion matrix
conf_matrix = confusion_matrix(y_true, y_pred)

# Plot confusion matrix heatmap
plt.figure(figsize=(10, 8))
disp = ConfusionMatrixDisplay(confusion_matrix=conf_matrix, display_labels=validation_generator.classes_)
disp.plot(cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.show()

# Print classification report
report = classification_report(y_true, y_pred, target_names=validation_generator.classes_)
print(report)

# Calculate and print metrics
accuracy = np.trace(conf_matrix) / np.sum(conf_matrix)
precision = np.diag(conf_matrix) / np.sum(conf_matrix, axis=0)
recall = np.diag(conf_matrix) / np.sum(conf_matrix, axis=1)
specificity = (np.sum(conf_matrix) - np.sum(conf_matrix, axis=0) - np.sum(conf_matrix, axis=1)) / (np.sum(conf_matrix) - np.sum(conf_matrix, axis=0) - np.sum(conf_matrix, axis=1))
f1_scores = 2 * (precision * recall) / (precision + recall)

print(f'Accuracy: {accuracy:.4f}')
print(f'Precision: {precision}')
print(f'Recall: {recall}')
print(f'Specificity: {specificity}')
print(f'F1-Scores: {f1_scores}')

# Evaluate DenseNet121 model
evaluate_model(densenet_model, validation_generator, steps_per_epoch_val)
```

1/1	5s	5s/step
1/1	1s	885ms/step
1/1	1s	883ms/step
1/1	1s	872ms/step
1/1	1s	887ms/step
1/1	1s	979ms/step
1/1	1s	963ms/step
1/1	1s	934ms/step
1/1	1s	930ms/step
1/1	1s	910ms/step
1/1	1s	918ms/step
1/1	1s	911ms/step
1/1	1s	930ms/step
1/1	1s	932ms/step
1/1	1s	937ms/step
1/1	1s	939ms/step
1/1	1s	980ms/step
1/1	1s	928ms/step
1/1	1s	930ms/step
1/1	1s	940ms/step
1/1	1s	923ms/step
1/1	1s	926ms/step
1/1	1s	950ms/step
1/1	1s	948ms/step
1/1	1s	904ms/step
1/1	1s	883ms/step
1/1	1s	890ms/step
1/1	1s	866ms/step
1/1	1s	874ms/step
1/1	1s	866ms/step
1/1	1s	855ms/step
1/1	1s	864ms/step
1/1	1s	858ms/step
1/1	1s	860ms/step
1/1	1s	871ms/step
1/1	1s	855ms/step
1/1	1s	860ms/step
1/1	1s	860ms/step
1/1	1s	860ms/step
1/1	1s	881ms/step
1/1	1s	873ms/step
1/1	1s	858ms/step
1/1	1s	863ms/step
1/1	1s	864ms/step
1/1	1s	861ms/step
1/1	1s	861ms/step
1/1	1s	865ms/step
1/1	1s	865ms/step
1/1	1s	860ms/step
1/1	1s	882ms/step
1/1	1s	866ms/step
1/1	1s	896ms/step
1/1	1s	870ms/step
1/1	1s	876ms/step
1/1	1s	855ms/step
1/1	1s	885ms/step


1/1	1s	865ms/step
1/1	1s	860ms/step
1/1	1s	860ms/step
1/1	1s	860ms/step
1/1	1s	859ms/step
1/1	1s	855ms/step
1/1	1s	936ms/step
1/1	1s	860ms/step
1/1	1s	864ms/step
1/1	1s	860ms/step
1/1	1s	886ms/step
1/1	1s	865ms/step
1/1	1s	883ms/step
1/1	1s	876ms/step
1/1	1s	870ms/step
1/1	1s	856ms/step
1/1	1s	847ms/step
1/1	1s	864ms/step
1/1	1s	850ms/step
1/1	1s	855ms/step
1/1	1s	850ms/step
1/1	1s	868ms/step
1/1	1s	850ms/step
1/1	1s	860ms/step
1/1	1s	855ms/step
1/1	1s	857ms/step
1/1	1s	860ms/step
1/1	1s	865ms/step
1/1	1s	860ms/step
1/1	1s	865ms/step
1/1	1s	867ms/step
1/1	1s	866ms/step
1/1	1s	869ms/step
1/1	1s	875ms/step
1/1	1s	868ms/step
1/1	1s	850ms/step
1/1	1s	870ms/step
1/1	1s	869ms/step
1/1	1s	853ms/step
1/1	1s	852ms/step
1/1	1s	865ms/step
1/1	1s	863ms/step
1/1	1s	868ms/step
1/1	1s	850ms/step
1/1	1s	883ms/step
1/1	1s	867ms/step
1/1	1s	865ms/step
1/1	1s	863ms/step
1/1	1s	855ms/step
1/1	1s	880ms/step
1/1	1s	865ms/step
1/1	1s	865ms/step
1/1	1s	850ms/step
1/1	1s	850ms/step
1/1	1s	850ms/step
1/1	1s	860ms/step

1/1	1s	840ms/step
1/1	1s	857ms/step
1/1	1s	849ms/step
1/1	1s	863ms/step
1/1	1s	840ms/step
1/1	1s	856ms/step
1/1	1s	854ms/step
1/1	1s	847ms/step
1/1	1s	845ms/step
1/1	1s	856ms/step
1/1	1s	843ms/step
1/1	1s	898ms/step
1/1	1s	878ms/step
1/1	1s	890ms/step
1/1	1s	862ms/step
1/1	1s	860ms/step
1/1	1s	860ms/step
1/1	1s	855ms/step
1/1	1s	869ms/step
1/1	1s	861ms/step
1/1	1s	848ms/step
1/1	1s	857ms/step
1/1	1s	852ms/step
1/1	1s	841ms/step
1/1	1s	850ms/step
1/1	1s	865ms/step
1/1	1s	853ms/step
1/1	1s	840ms/step
1/1	1s	845ms/step
1/1	1s	866ms/step
1/1	1s	859ms/step
1/1	1s	866ms/step
1/1	1s	850ms/step
1/1	1s	837ms/step
1/1	1s	860ms/step
1/1	1s	851ms/step
1/1	1s	842ms/step
1/1	1s	850ms/step
1/1	1s	840ms/step
1/1	1s	850ms/step
1/1	1s	860ms/step
1/1	1s	840ms/step
1/1	1s	862ms/step
1/1	1s	844ms/step
1/1	1s	868ms/step
1/1	1s	883ms/step
1/1	1s	869ms/step
1/1	1s	904ms/step
1/1	1s	875ms/step
1/1	1s	886ms/step
1/1	1s	851ms/step
1/1	1s	840ms/step
1/1	1s	850ms/step
1/1	1s	846ms/step
1/1	1s	855ms/step
1/1	1s	844ms/step

1/1	1s	850ms/step
1/1	1s	838ms/step
1/1	1s	857ms/step
1/1	1s	855ms/step
1/1	1s	875ms/step
1/1	1s	850ms/step
1/1	1s	868ms/step
1/1	1s	845ms/step
1/1	1s	859ms/step
1/1	1s	850ms/step
1/1	1s	848ms/step
1/1	1s	853ms/step
1/1	1s	840ms/step
1/1	1s	856ms/step
1/1	1s	847ms/step
1/1	1s	849ms/step
1/1	1s	884ms/step
1/1	1s	842ms/step
1/1	1s	830ms/step
1/1	1s	853ms/step
1/1	1s	880ms/step
1/1	1s	839ms/step
1/1	1s	851ms/step
1/1	1s	849ms/step
1/1	1s	841ms/step
1/1	1s	845ms/step
1/1	1s	833ms/step
1/1	1s	838ms/step
1/1	1s	845ms/step
1/1	1s	850ms/step
1/1	1s	860ms/step
1/1	1s	862ms/step
1/1	1s	864ms/step
1/1	1s	840ms/step
1/1	1s	840ms/step
1/1	1s	840ms/step
1/1	1s	847ms/step
1/1	1s	850ms/step
1/1	1s	859ms/step
1/1	1s	870ms/step
1/1	1s	977ms/step
1/1	1s	917ms/step
1/1	1s	895ms/step
1/1	1s	906ms/step
1/1	1s	920ms/step
1/1	1s	916ms/step
1/1	1s	911ms/step
1/1	1s	925ms/step
1/1	1s	965ms/step
1/1	1s	1s/step
1/1	1s	904ms/step
1/1	1s	865ms/step
1/1	1s	853ms/step
1/1	1s	843ms/step
1/1	1s	843ms/step
1/1	1s	843ms/step

1/1	1s	856ms/step
1/1	1s	852ms/step
1/1	1s	870ms/step
1/1	1s	864ms/step
1/1	1s	858ms/step
1/1	1s	858ms/step
1/1	1s	863ms/step
1/1	1s	874ms/step
1/1	1s	870ms/step
1/1	1s	891ms/step
1/1	1s	863ms/step
1/1	1s	913ms/step
1/1	1s	850ms/step
1/1	1s	865ms/step
1/1	1s	892ms/step
1/1	1s	841ms/step
1/1	1s	858ms/step
1/1	1s	857ms/step
1/1	1s	846ms/step
1/1	1s	850ms/step
1/1	1s	847ms/step
1/1	1s	855ms/step
1/1	1s	845ms/step
1/1	1s	853ms/step
1/1	1s	844ms/step
1/1	1s	842ms/step
1/1	1s	865ms/step
1/1	1s	845ms/step
1/1	1s	846ms/step
1/1	1s	845ms/step
1/1	1s	840ms/step
1/1	1s	832ms/step
1/1	1s	845ms/step
1/1	1s	863ms/step
1/1	1s	863ms/step
1/1	4s	4s/step
1/1	1s	861ms/step
1/1	1s	850ms/step
1/1	1s	864ms/step
1/1	1s	864ms/step

<Figure size 1000x800 with 0 Axes>



	precision	recall	f1-score	support
COVID	0.94	0.88	0.91	1431
Lung_Opacity	0.90	0.89	0.89	2403
Normal	0.91	0.95	0.93	4065
Viral Pneumonia	0.95	0.93	0.94	533
accuracy			0.92	8432
macro avg	0.93	0.91	0.92	8432
weighted avg	0.92	0.92	0.92	8432

Accuracy: 0.9164
Precision: [0.94011976 0.90245971 0.91208531 0.95366795]
Recall: [0.8777079 0.88555972 0.94686347 0.92682927]
Specificity: [0.98857306 0.96185105 0.91504465 0.99696164]
F1-Scores: [0.90784243 0.89392985 0.92914906 0.94005709]

In []: